

Lower bounding the description complexity of quantum states

Chinmay Nirkhe (IBM Quantum Cambridge)

based on joint works with

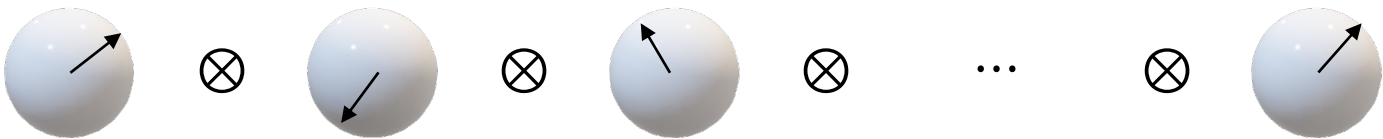
[1] A. Natarajan (MIT),

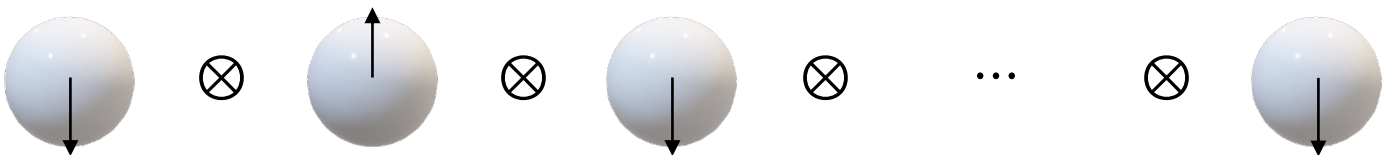
[2] A. Anshu (Harvard), N. Breuckmann (Bristol),

[3] S. Irani (Simons Inst.), A. Natarajan (MIT), S. Rao (MIT), H. Yuen (Columbia)

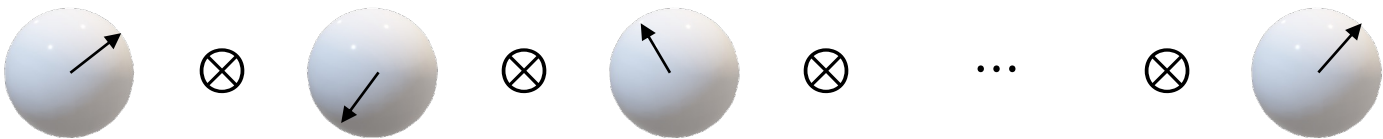


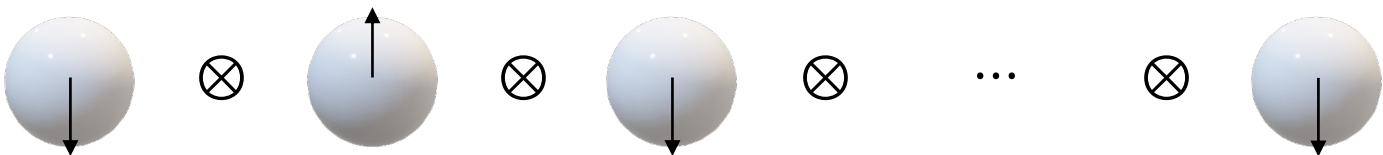
We know that most quantum states are complex...

Quantum:  $\approx 2^{2^n}$ states

Classical:  $= 2^n$ states

but how many of
them are interesting
for physics?

Quantum:  $\approx 2^{2^n}$ states


Classical:  $= 2^n$ states

Quantifying the states that are interesting for physics

The energy operator in quantum mechanics is called the **Hamiltonian**.

Interesting physical systems are defined by Hamiltonians of a special form called **local Hamiltonians**.

$$\mathbf{H} = \sum h_i$$



$$h_i = |000\rangle\langle 000| - |111\rangle\langle 111|$$

Due to the importance of ground states in condensed matter physics,

we would really like to know if ground states of local Hamiltonians have efficient (i.e. short) and verifiable classical descriptions.

This talk:

We prove lower bounds on the classical description complexity of ground (and low-energy) states of local Hamiltonians.

Are there local Hamiltonian ground states that cannot be described by polynomial depth quantum circuits?

The converse is false [Kitaev⁰³].
For every circuit \mathcal{C} , there exists
a LH with ground state \approx
 $\mathcal{C}|0^n\rangle \otimes |\text{junk}\rangle$.


A counting exercise ...

How many local Hamiltonians are there?

Each h_i can be described by $O(2^{2\ell} \log n)$ bits.

There are n^ℓ terms, so \mathbf{H} can be described by $O(2^{2\ell} n^\ell \log n)$ bits.

So, there are at most $2^{\text{poly}(n)}$ local Hamiltonians for $\ell = O(1)$.



$$h_i = |000\rangle\langle 000| - |111\rangle\langle 111|$$


How many depth t quantum circuits are there?

Each row of the quantum circuit can be described by a matching and $n/2$ 4×4 matrices.

So, there are $2^{O(t n^2)}$ quantum circuits of depth t .

Therefore, counting arguments alone don't show separation and, secondly, both the number of local Hamiltonians and "accessible" quantum states are \ll number of quantum states (2^{2^n}).

The space of quantum states



Outline of today's talk

Oracle separations between QCMA and QMA

*i.e., proving polynomial lower bounds on describing ground states
(for a generalization of local Hamiltonians)*

The NLTS problem

*i.e., proving logarithmic-depth circuit lower bounds for
low-energy states of local Hamiltonians*

Quantum search-to-decision reductions

*i.e., why we cannot study the complexity of quantum states via
the study of quantum decision problems*


A classical oracle
separations for
QCMA and QMA

The connection to quantum complexity theory

The energy operator in quantum mechanics is called the **Hamiltonian**.

Interesting physical systems are defined by Hamiltonians of a special form called **local Hamiltonians**.

$$\mathbf{H} = \sum h_i$$



$$h_i = |000\rangle\langle 000| - |111\rangle\langle 111|$$

The problem of calculating the ground energy

$$E = \min_{|\psi\rangle} \langle \psi | \mathbf{H} | \psi \rangle$$

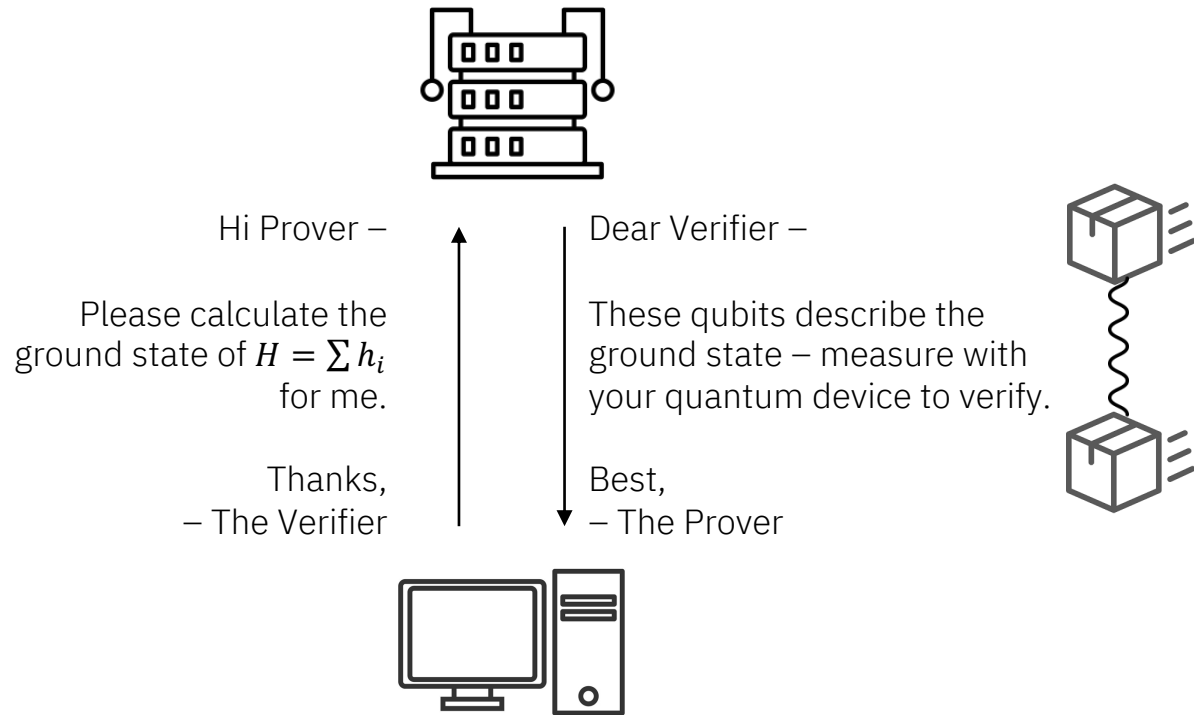
famously connects physics to computer science as the problem is *complete* for the class **QMA**.

QMA = Quantum Merlin-Arthur

If the ground state can always be *verifiably* classically described, then the problem is complete for the class **QCMA**.

QCMA = Quantum-Classical Merlin-Arthur

QMA



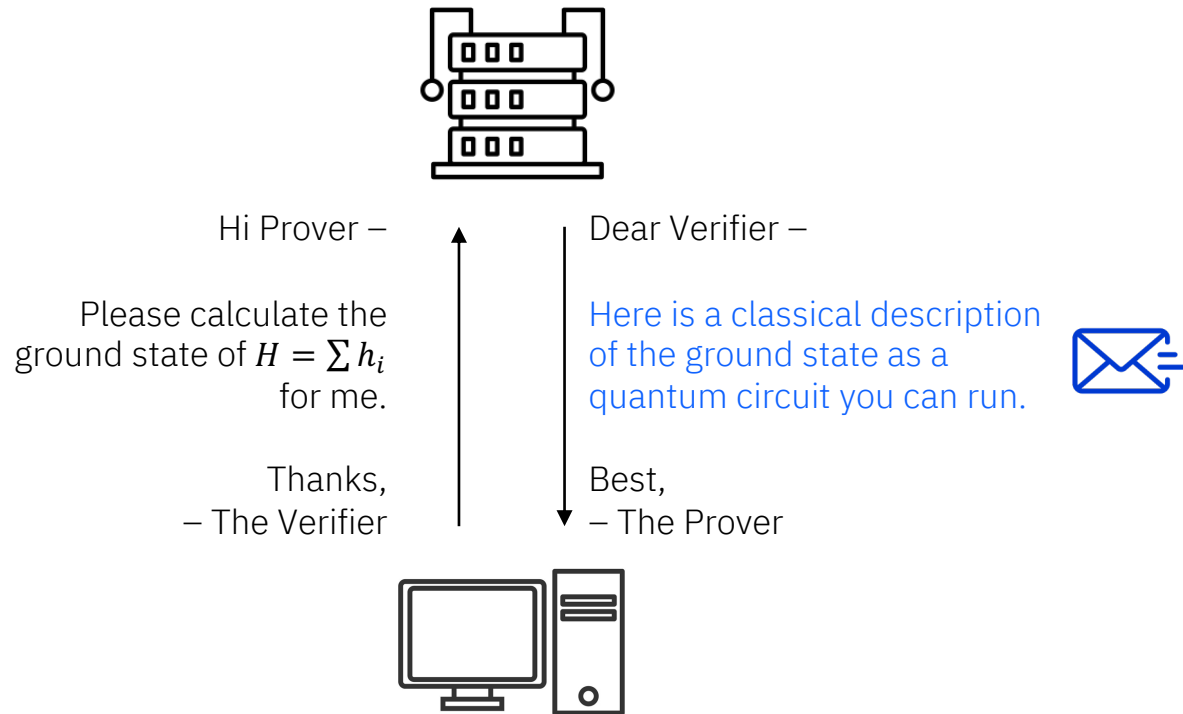
QMA: The set of problems that can be verified by a prover sending a quantum state to the verifier which the verifier uses in a quantum circuit.

Complete problem: Deciding if there exists a quantum state $|\psi\rangle$ such that

YES: $\langle\psi|H|\psi\rangle \leq a$

NO: or else, $\langle\psi|H|\psi\rangle \geq a + 1/n^3$ for all states

QCMA



QCMA: The set of problems that can be verified by a prover sending a **classical string** to the verifier which the verifier uses in a quantum circuit.

Complete problem: Deciding if there exists a quantum state $|\psi\rangle$ such that

YES: $\langle\psi|H|\psi\rangle \leq a$

NO: or else, $\langle\psi|H|\psi\rangle \geq a + 1/n^3$ for all states

where in both cases, $|\psi\rangle$ is the output of a quantum circuit of polynomial depth.


The question, of whether QCMA and QMA are equal, is a major open question in quantum complexity theory.

QCMA \neq QMA


implies

that not all local Hamiltonians
ground states have short verifiable
classical descriptions.

Contrapositive: If short descriptions *always* exist, then prover can
always send the short description instead of the quantum state.



vs.



DALL-E 2 renderings.

Theorem [Natarajan-Nirkhe²²]: There is a black-box distribution D for which we can prove that


$$\text{QCMA}^D \neq \text{QMA}^D.$$

This is the strongest evidence yet that ground states cannot be classically described.

Proving $\text{QCMA} \neq \text{QMA}$ outright would have incredible implications for complexity theory. Consequently, it requires truly novel techniques.




Instead, the best we can do is slightly generalize the notion of local Hamiltonians until we can prove $\text{QCMA} \neq \text{QMA}$.




Generalizing sparse Hamiltonians to graphs

$$\mathbf{H} = \begin{pmatrix} 0 & \alpha & 0 & 0 & 0 & \theta & 0 & 0 \\ \alpha & 0 & 0 & 0 & 0 & \epsilon & 0 & 0 \\ 0 & 0 & 0 & \beta & 0 & 0 & \delta & 0 \\ 0 & 0 & \beta & \lambda & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & \gamma & 0 & 0 & \iota \\ \theta & \epsilon & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & \delta & 0 & 0 & 0 & 0 & \zeta \\ 0 & 0 & 0 & 0 & \iota & 0 & \zeta & 0 \end{pmatrix}$$

n -qubit sparse Hamiltonian



2^n vertex weighted graph



2^n vertex unweighted graph



$$\mathbf{H} = \begin{pmatrix} 0 & 1 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 \end{pmatrix}$$

Adjacency matrix of graph

Low energies of sparse graph Hamiltonians

If \mathbf{H} is the Hamiltonian corresponding to a d -regular sparse graph,

then $-\mathbf{H}$ is a Hamiltonian with ground energy equal to $-d$ with a ground state of $\sum_{x \in \{0,1\}^n} |x\rangle$, the uniform super position.

What is the second smallest eigenvalue?

If the graph has ≥ 2 connected components, it is also $-d$.

If the graph is α -expanding, it is $-d + \alpha d$.


Theorem 1: Deciding if a graph (given as black-box sparse adjacency list) either has multiple connected components or is 0.01 -expanding is in QMA.

This also holds for certain distributions over graphs.

Theorem 2: For the same set of distributions, we can show that no efficient QCMA algorithm exists.

Together they prove $\text{QCMA}^D \neq \text{QMA}^D$.

The **QMA** algorithm



$$|\xi_1\rangle = \frac{|3\rangle + |4\rangle + |5\rangle + |7\rangle + |8\rangle}{\sqrt{5}}$$

$$|\xi_2\rangle = \frac{|1\rangle + |2\rangle + |6\rangle}{\sqrt{3}}$$

Easy to check that

$$|\xi_{\text{solution}}\rangle = \frac{\sqrt{5}}{\sqrt{8}}|\xi_2\rangle - \frac{\sqrt{3}}{\sqrt{8}}|\xi_1\rangle$$

is a eigenvector of eigenvalue $-d$ as well and is orthogonal to $\sum_{x \in \{0,1\}^n} |x\rangle$.


Quantum solution is to provide $|\xi_{\text{solution}}\rangle$ which proves that there are 2 eigenvectors of eigenvalue $-d$.

If the graph is α -expanding, this test **fails** with probability $\alpha/4$ as there is only 1 eigenvector of eigenvalue $-d$ and the next eigenvalue is $-d + \alpha d$.

Sketch of **QCMA** impossibility result

In the **QMA** algorithm, we saw that the solution state only depends on the *vertices* in the connected component.


Let's first try to prove an easier subproblem: an impossibility result for all classical proofs that only depend on the *vertices* in one connected component.




Notation: $N = 2^n$ (number of vertices).
 $z = N^{9/10}$.

Setup: YES instance graphs are promised to have N/z expanding connected components each of size z . NO instance graphs have 1 expanding component.

Assume that for every set $S \subset [N]$ of size z , there exists exactly one YES instance that has S has a connected component and that the proof for said YES instance depends only on S .



YES instance



NO instance

Many YES instances correspond to the same proof π

Assume there exists a **bijection** between [YES instance graphs] and [sets of size z].

Then, there are at least $\binom{N}{z} \geq 2^{N^{0.9}}$ YES instances.

And, the number of proofs of length q is at most 2^q .

Pidgeon-hole principle: If there was a QCMA algorithm for all YES instances with proof of length q , then there exists a popular proof $\pi \in \{0,1\}^q$ such that at least $2^{N^{0.9}-q} \geq 2^{N^{0.8}}$ graphs correspond to π .

By the **bijection**, $2^{N^{0.8}}$ [sets of size z] correspond to π .


Ramsey theory: “If the set is large enough, there exists a subset which is structured”

Within the unstructured set of YES instances that correspond to π , there exists a set of YES instances that are structured: they form a sunflower:

Sunflower: A collection Σ of sets of size z such that

- (1) (core) $\exists F \subset [N]$ such that $\forall S \in \Sigma, F \subset S$.
- (2) $\forall x \in (\bigcup_{S \in \Sigma} S) \setminus F$, then

$$\Pr_{S \in \Sigma}[x \in S] \leq \left(\frac{z}{N}\right)^{0.99}.$$



Fixing the proof of a **QCMA** algorithm yields a **BQP** algorithm

Consider the **BQP** algorithm \mathcal{A} generated by fixing the popular proof π into the **QCMA** algorithm.

- \mathcal{A} must answer YES (with high probability) on every graph corresponding to a subset S in the sunflower Σ .
- \mathcal{A} must answer NO (with high probability) on every NO instance graph.

We will show that any algorithm \mathcal{A} achieving these two tasks requires exponential time.

Ramsey theory: “If the set is large enough, there exists a subset which is structured”


Within the unstructured set of YES instances that correspond to π , there exists a set of YES instances that are structured: they form a sunflower:

Sunflower: A collection Σ of sets of size z such that

(1) (core) $\exists F \subset [N]$ such that $\forall S \in \Sigma, F \subset S$.

(2) $\forall x \in (\bigcup_{S \in \Sigma} S) \setminus F$, then

$$\Pr_{S \in \Sigma}[x \in S] \leq \left(\frac{z}{N}\right)^{0.99}.$$



The adversary method


\mathcal{A} must answer YES (with high probability) on every graph corresponding to a subset S in the sunflower Σ .

Sunflower: A collection Σ of sets of size z such that

(1) (core) $\exists F \subset [N]$ such that $\forall S \in \Sigma, F \subset S$.

(2) $\forall x \in (\bigcup_{S \in \Sigma} S) \setminus F$, then

$$\Pr_{S \in \Sigma}[x \in S] \leq \left(\frac{z}{N}\right)^{0.99}.$$



Ideal sunflower corresponding to F :


$$\Sigma_{\text{ideal}} = \{S \subset [N] \mid |S| = z, F \subset S\}.$$

Then by assumption, \mathcal{A} must answer YES (with high probability) on every graph corresponding to a subset S in the sunflower Σ_{ideal} .

Lemma: Any quantum algorithm cannot distinguish a graph sampled from Σ and a graph sampled from Σ_{ideal} without making an exponential number of queries to the graph.

Random walk sampling reduction

\mathcal{A} must answer YES (with high probability) on every graph corresponding to a subset S in the sunflower Σ_{ideal} .



We now generate an algorithm \mathcal{A}' which has F hard-coded. \mathcal{A}' will not be time-efficient (just query-efficient).

Starting with a random point v_1 , \mathcal{A}' applies a random walk sampling to sample points

$$V = \{v_1, v_2, \dots, v_{|F|}\}$$

from the same connected component.


\mathcal{A}' picks a permutation $\sigma: [N] \rightarrow [N]$ mapping $V \rightarrow F$.

Then, \mathcal{A}' simulates \mathcal{A} except when a query to vertex w is made, \mathcal{A}' queries $\sigma^{-1}(w)$.

By the correctness of \mathcal{A} and the random walk being a good sampler (our YES instances are expanding within components), \mathcal{A}' must also succeed with high probability.

This gives a **BQP** algorithm \mathcal{A}' for deciding the original problem (modulo a few technicalities).

Appeal to known quantum query lower bounds



[Ambainis-Childs-Liu¹¹] proved a quantum query lower bound for the problem (without proof).

Overall, we have shown that any *classical proof* depending only on the vertices in one connected component cannot help with this problem.

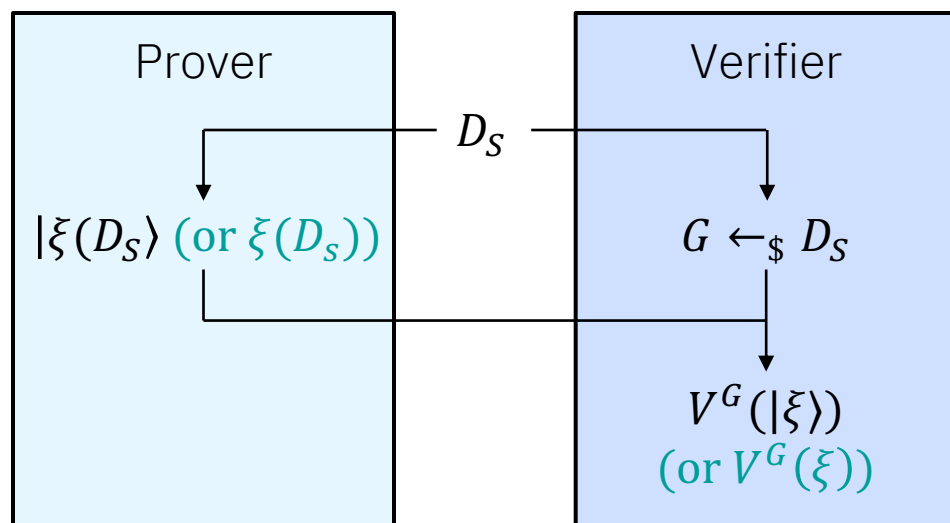
In other words, a “*near-sighted*” classical prover – **one who can only see the overall structure of the graph (connected components) but not internal structures (such as edges, triangles, etc.)** – cannot help.

But recall from our **QMA** algorithm, a “near-sighted” quantum prover can help!

Can we “blind” the classical prover so that he is near-sighted?

Blinding the classical prover

Let D_S be the uniform distribution over all graphs that have a connected component S for S of size z .



Meaning, the prover only knows the graph up to the distribution. This is equivalent to the prover only answering based on the connected component S – this is the same as blinding the prover until it is "near-sighted".

Blinding the classical prover yields our impossibility result while blinding the quantum prover has no effect.

This proves that there is a *distribution over graphs* which separates **QMA** and **QCMA**.

Removing the "blinding" step seems difficult as there are known obstacles that need avoiding.

Theorem [Natarajan-Nirkhe²²]: There is a black-box distribution D for which we can prove that


$$\text{QCMA}^D \neq \text{QMA}^D.$$

This is the strongest evidence yet that ground states cannot be classically described.

The NLTS problem
and the QPCP
conjecture

The space of quantum states

and the complexity class that they serve as witness for



The hardness of low-energy states


QPCP Conjecture: For some $\epsilon > 0$, it is QMA-hard to decide whether the ground energy of a local Hamiltonian $\mathbf{H} = \sum_{i=1}^m h_i$ is

(yes) small: $\langle \psi | H | \psi \rangle \leq a$

(no) or large: $\langle \psi | H | \psi \rangle \geq a + \epsilon m$

Simple consequence: Any state $|\psi'\rangle$ of energy $\leq a + \frac{\epsilon}{2}m$ is also a verifiable if the QPCP conj. is true

Therefore, the set of witnesses not only includes ground states but also low-energy states



NLTS Thm [Anshu-Breuckmann-Nirkhe²²]: There exists a local Hamiltonian s.t. no low-energy states ($\leq \epsilon n$) is the output of a constant-depth circuit.

QPCP Conj. + $\text{NP} \neq \text{QMA} \Rightarrow \text{NLTS}$

i.e., a necessary, but not sufficient, consequence

NLTS is an unconditional statement, unlike the previous **QCMA v QMA** separation.


Sketch of the **NLTS** proof

We first need a technique for proving logarithm-depth circuit depth lower bounds.

Local indistinguishability: Two states $|\psi\rangle$ and $|\psi'\rangle$ are d -locally indistinguishable if every reduced density matrix on d qubits is the same.


Fact: d -locally indistinguishable states have $\Omega(\log d)$ circuit-depth lower bounds.

Corollary: If measuring a quantum state yields a probability distribution p such that there exist two subsets $S_1, S_2 \subset \{0,1\}^n$ with $p(S_1), p(S_2) \geq \Omega(1)$ and the Hamming distance between S_1 and S_2 is $\Omega(n)$, then the quantum state has a $\Omega(\log n)$ circuit depth lower bound.



The challenge is to find a local Hamiltonian for which every low-energy state when measured yields such a “well-spread” distribution.

The low-energy subset of a linear-distance error-correcting code is supported like a “well-spread” dist.




Sketch of the **NLTS** proof

By considering quantum LDPC error-correcting codes of constant-rate and linear-distance (plus an additionally robustness property), we can find the “well-spread” property and construct local Hamiltonians with the **NLTS** property.


The key is using the *uncertainty lemma* to argue that for any state, measuring in either the standard or Hadamard basis must yield a well-spread distribution.

This is because the quantum code simultaneously corrects large X- and Z-errors.



The challenge is to find a local Hamiltonian for which every low-energy state when measured yields such a “well-spread” state.

The low-energy subset of a linear-distance error-correcting code is supported like a “well-spread” dist.




Proving stronger lower bounds than **NLTS**

Constant-depth quantum circuits are just one of *many* classical witness that can be provided for an **NP** proof.

QPCP Conj. + $\mathbf{NP} \neq \mathbf{QMA} \Rightarrow$
lower-bounds for all families of **NP** witnesses

Open question: Can we prove lower-bounds for some other families of **NP** witnesses? Is there is a family of local Hamiltonians for which all known **NP** witnesses are insufficient?



Any state of this form is also a **NP** witness.

NLTS+ conjecture: There exists local Hamiltonian such that all such states have energy $\geq \epsilon n$.

Our proof of **NLTS** does not satisfy this!

Quantum search-
to-decision
reductions

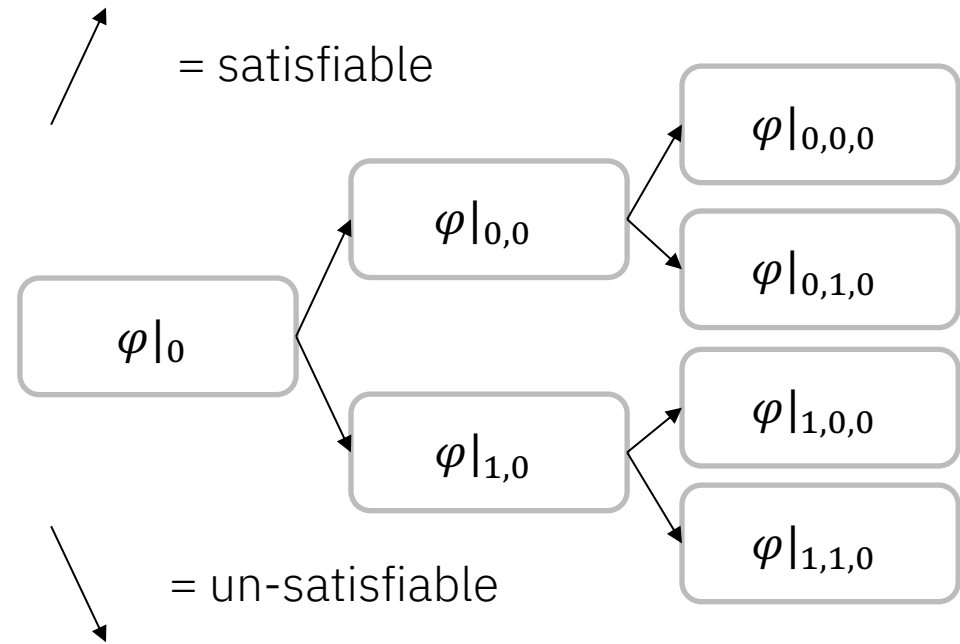
NP has search-to-decision reductions

Say there is a black box which takes as input 3SAT formulas φ and outputs (with $\text{pr} = 1$) if they are satisfiable or not...

Crucially, it does not tell you the solution (satisfying assignment) $x \in \{0,1\}^n$ s.t. $\varphi(x) = 1$.

This is a black box for NP and repeated uses can be used to build a solution $x \in \{0,1\}^n$.

Let $\varphi|_{a_1, a_2, \dots, a_k}$ be the restriction of φ on the first k variables.



After n queries, we learn a complete satisfying assignment.

$$\text{NP-Search} \subseteq \text{P}^{\text{Decision-NP}}$$

Does **QMA** have search-to-decision reductions?

In classical CS theory, defining decision problems as the *de facto* model of computation is justified by search-to-decision reductions.

What about in quantum CS? Is the same definition justified? Or do we need to rectify our *de facto* notion of computation?

[Irani-Natarajan-Nirkhe-Rao-Yuen²¹]:

Theorem 1: QMA-search is reducible to 1-query PP-decision

Theorem 2: Oracle proof that QMA-search not reducible to QMA-decision

Theorem 1: Let (\mathbf{H}, a, b) be a local Hamiltonian problem. If the problem is a YES instance, there exists a BQP algorithm making 1-query to a PP-oracle such that the output state of the algorithm has energy $\leq a + (b - a)/2$.

Theorem 2: We show that a unitary oracle corresponding to a QCMA v QMA separation also proves that QMA-search is not reducible to QMA-decision

It's open whether our distribution oracle also proves such an impossibility result.

Understanding quantum states

The difference in quantum search-to-decision reductions and classical search-to-decision reductions suggests that quantum states cannot be entirely studied through the lens of decision problems.

We need to better understand the complexity of quantum states and ideally how to prove polynomial-depth lower bounds for quantum states without oracles.

Other open questions:

How can we remove the distributions from the oracle separations between **QCMA** and **QMA**?

Can we use the techniques to prove oracle separations for other quantum complexity classes such as **QMA(2)** and **QMA**?

What is the power of **BQP^{QMA}**? It lies somewhere between **QMA** and **QCMA**.

Thank you for listening.

Chinmay Nirkhe (IBM Quantum Cambridge)

